



UNIVERSITY OF MARYLAND BALTIMORE COUNTY

DEPARTMENT OF COMPUTER SCIENCE

ADVANCED OPERATING SYSTEMS

CMSC 621

PROJECT 3

SUBMITTED BY:

SUSHMITHA MANJUNATHA

ABSTRACT

In this project, I have tried to implement a distributed banking service that permits multiple concurrent operations and exhibits simple fault tolerance. You should run 3 servers, each as a separate process on your machine. Each server should be capable of handling multiple simultaneous requests, and any account should be accessible from any server whether or not it was created on that server. Additionally, the group of servers should tolerate the loss of any one server in the group without impacting the overall service. You can kill a process to emulate a crash. What this means is that you must replicate the data between servers. Servers should maintain just enough state as to complete the transactions. For locking, you can use any locking mechanism. For atomicity, you should use the two-phase commit protocol. Write your server and client in C++.

SYSTEM DESIGN

Description of the functions used:

In-built functions used:

argc - It is the number of arguments passed into the program from the command line, including the name of the program.

argv[] - The array of character pointers is the listing of all the arguments provided in the command line.

stderr - Standard error is an output stream typically used by programs to output error messages or diagnostics. It is a stream independent of standard output and can be redirected separately.

gethostbyname() – This function returns a structure of type *hostent* for the given host *name*. Here *name* is either a hostname or an IPv4 address in standard dot notation.

atoi - Convert string to integer.

atof - Convert string to double.

sockaddr_in - The basic structure for all syscalls and functions that deal with internet addresses.

socket(domain,type,protocol) - The function socket() creates an endpoint for communication and returns a file descriptor for the socket. socket() takes three arguments: domain, type and protocol.

AF_INET- Indicates network protocol IPv4 (IPv4-only).

SOCK_STREAM – It is reliable stream-oriented service .

bind() - It assigns a socket to an address. When a socket is created using socket(), it is only given a protocol family, but not assigned an address. This association with an address must be performed with the bind() system call before the socket can accept connections to other hosts. bind() takes three arguments:

sockfd, a descriptor representing the socket to perform the bind on.

`my_addr`, a pointer to a `sockaddr` structure representing the address to bind to.

`addrlen`, a `socklen_t` field specifying the size of the `sockaddr` structure.

`Bind()` returns 0 on success and -1 if an error occurs.

listen() – After a socket has been associated with an address, `listen()` prepares it for incoming connections. However, this is only necessary for the stream-oriented (connection-oriented) data modes, i.e., for socket types (`SOCK_STREAM`, `SOCK_SEQPACKET`). `listen()` requires two arguments:

`sockfd`, a valid socket descriptor.

`backlog`, an integer representing the number of pending connections that can be queued up at any one time. The operating system usually places a cap on this value.

Once a connection is accepted, it is dequeued. On success, 0 is returned. If an error occurs, -1 is returned.

ifstream - Stream class to read from files

fstream - Stream class to both read and write from/to files.

htons() - This function converts the unsigned short integer *hostshort* from host byte order to network byte order.

filename.c_str()- If *filename* is a variable of type *string*, we can obtain the corresponding C-style string by calling the function *filename.c_str()*.

bzero(buff,n) - The `bzero()` function erases the data in the *n* bytes of the memory starting at the location pointed to by *buff*, by writing zeroes (bytes containing '\0') to that area.

sprint(char *str ,const char *format) - Composes a string with the same text that would be printed if *format* was used on printf, but instead of being printed, the content is stored as a C string in the buffer pointed by *str*.

write(int fd, const void *buf, size_t count) - It writes up to *count* bytes from the buffer pointed *buf* to the file referred to by the file descriptor *fd*.

read(int fd, void *buf, size_t count) - attempts to read up to *count* bytes from file descriptor *fd* into the buffer starting at *buf*.

sockfd - It is the socket descriptor returned by `socket()`. `serv_addr` is pointer to struct `sockaddr` that contains information on destination IP address and port. `addrlen` is set to `sizeof(struct sockaddr)`

sleep() – It makes the calling thread sleep until specified seconds have elapsed or a signal arrives which is not ignored.

accept() - When an application is listening for stream-oriented connections from other hosts, it is notified of such events and must initialize the connection using the `accept()` function. The `accept()` function creates a new socket for each connection and removes the connection from the listen queue. It takes the following arguments:

`sockfd`, the descriptor of the listening socket that has the connection queued.

`cliaddr`, a pointer to a `sockaddr` structure to receive the client's address information.

`addrlen`, a pointer to a `socklen_t` location that specifies the size of the client address structure passed to `accept()`. When `accept()` returns, this location indicates how many bytes of the structure were actually used.

The `accept()` function returns the new socket descriptor for the accepted connection, or -1 if an error occurs.

pthread_attr_init(&tattr)- This function initializes the thread attributes object pointed to by *attr* with default attribute values.

pthread_attr_setdetachstate((&tattr, int detachstate) – This function sets the detach state attribute of the thread attributes object referred to by *tattr* to the value specified in *detachstate*. The detach state attribute determines whether a thread created using the thread attributes object *tattr* will be created in a joinable or a detached state.

pthread_create() - This function starts a new thread in the calling process. The new thread starts execution by invoking *start_routine()*; *arg* is passed as the sole argument of *start_routine()*.

pthread_mutex_init(&mutex, NULL) - This function shall initialize the mutex referenced by *mutex* with attributes specified by *attr*. If *attr* is NULL, the default mutex attributes are used; the effect shall be the same as passing the address of a default mutex attributes object. Upon successful initialization, the state of the mutex becomes initialized and unlocked.

pthread_mutex_lock(&mutex) - The mutex object referenced by *mutex* is locked by calling *pthread_mutex_lock()*. If the mutex is already locked, the calling thread blocks until the mutex becomes available. This operation returns with the mutex object referenced by *mutex* in the locked state with the calling thread as its owner.

pthread_mutex_unlock(&mutex) - This function shall release the mutex object referenced by *mutex*. The manner in which a mutex is released is dependent upon the mutex's type attribute.

Structures Used:

records - It stores the records of the accounts in the bank-server. It has attributes like account-number and balance.

How to run the code:

1. `g++ backendS.cpp -o backendS -lpthread`

2. `g++ frontendS.cpp -o frontendS -lpthread`

3. `g++ clientS.cpp -o clientS`

Backend server:

4. `./backendS 9090 1`

5. `./backendS 9091 2`

6. `./backendS 9092 3`

FrontendServer side

7. `./frontend localhost 8888 9090 9091 9092`

Client:

8. `./clientS <localhost> <clientportnumber>`

Frontend server

1. Three separate buffers are being used for each of the 3 servers.
2. Accepts 6 command line arguments i.e client_port_number, and port_numbers of all 3 servers along with the file name.
3. flag "NW" is being used to keep track of number of not working servers initially it will be 0 assuming all servers work.
4. (setsockopt(socketfd, SOL_SOCKET, SO_REUSEADDR, (char *) &reuse, sizeof(reuse)) - this function is being used SO_REUSEADDR Specifies that the rules used in validating addresses supplied to bind() should allow reuse of local addresses, if this is supported by the protocol. This option takes an int value. This is a Boolean option.
5. Now socket connection is made with 3 servers.
6. Connection is made with the client process
7. pthread_create(&t , NULL , handler_function , (void*) new_socketfd) < 0 - is used MULTITHREADING is done for handling each client request.
8. handler_function() - handles 2phase commit and locking scheme
9. reads multiple requests from clients
10. **pthread_mutex_lock(&mutex)- mutexs** are used for locking
11. Timeout values are set using "struct timeval"
11. Based on value of NW flag, COMMAND will be sent to the server

12.**(setsockopt (socketfd1, SOL_SOCKET, SO_RCVTIMEO, (char *)&timeout_values,sizeof(struct timeval))**-- is to use setsockopt() to set the socket's SO_RCVTIMEO option.

13.OK message is being read from all 3 servers.

14.when OK has been received from all the servers check if response has been received from all of them if not, then think of maintaining data consistency

15.If all servers are ready to commit,send global commit to all of them.

16.If even one of them sends a NO, ABORT the transaction and restart the transaction.

17.if one of them do not reply/ one server crashes COMMIT is sent

18.Read OK from Servers.

19.Write the response to the client.

Backend Servers

1.Structure called RECORDS is used for storing the account_number and account_balance.

2.It takes 3 command line arguments i.e servername, coordinator_port and Server_number

3.'Transaction_record' file will be created for each of the server

4.fopen(file_name,"w")- is used to read the file with write permission

5.**setsockopt(socketfd, SOL_SOCKET, SO_REUSEADDR, (char *) &reuse, sizeof(reuse))**- used to reuse the port

6.listening to connect to the clients

7. Reading the shared buffer for the commands from frontend server
8. copy the transaction to "transaction" array
9. Writing YES or NO to the coordinator.
10. reading commit/abort msg from coordinator
11. check if it is global commit msg or abort msg
12. parse transaction array and perform the task
13. If create , create an account id and return the account number
14. If QUERY , query for checking account balance and return account number
15. If UPDATE, update a record and return account number
16. If QUIT, send ready to commit msg to coordinator
17. If it is global abort dont save data into file

CLIENT functionality

1. Socket structures are defined
2. Client accepts hostname and portnumber as the argument
3. Connects to the frontend server
4. Receives OK from the coordinator
5. Receives commands from the user
6. writing the command input taken by user to buffer
7. converting the command to uppercase if the command is in lowercase
by **using islower() and toupper()** functions.
8. writing to the coordinator

9.reading the final response from the coordinator

10.checking whether it is a OK msg from coordinator to close the connection.

11.closing the socket

TWO PHASE COMMIT PROTOCOL

Commit request phase

or **voting phase**

1. The coordinator sends a **query to commit** message to all cohorts and waits until it has received a reply from all cohorts.
2. The cohorts execute the transaction up to the point where they will be asked to commit. They each write an entry to their *undo log* and an entry to their [*redo log*](#).
3. Each cohort replies with an **agreement** message (cohort votes **Yes** to commit), if the cohort's actions succeeded, or an **abort** message (cohort votes **No**, not to commit), if the cohort experiences a failure that will make it impossible to commit.

Commit phase

or **Completion phase**

Success[

If the coordinator received an agreement message from *all* cohorts during the commit-request phase:

1. The coordinator sends a **commit** message to all the cohorts.
2. Each cohort completes the operation, and releases all the locks and resources held during the transaction.
3. Each cohort sends an **acknowledgment** to the coordinator.
4. The coordinator completes the transaction when all acknowledgments have been received.

Failure

If *any* cohort votes **No** during the commit-request phase (or the coordinator's timeout **expires**):

1. The coordinator sends a **rollback** message to all the cohorts.
2. Each cohort undoes the transaction using the undo log, and releases the resources and locks held during the transaction.
3. Each cohort sends an **acknowledgement** to the coordinator.
4. The coordinator undoes the transaction when all acknowledgements have been received.

OUTPUT :

itna@shusmitha-VirtualBox: ~

```
shusmitha@shusmitha-VirtualBox:~$ make
g++ backends.cpp -o backends -lpthread
g++ frontendS.cpp -o frontendS -lpthread
g++ clients.cpp -o clients
```

```
shusmitha@shusmitha-VirtualBox:~$ ./backends 9090 1
Transaction_record1.txt
1: in server:CREATE 80
```

```
1 : commit:COMMIT
Going to commit
```

```
CREATE
```

```
buffer:OK 100
```

```
1: in server:QUERY 100
```

```
1 : commit:COMMIT
Going to commit
```

```
QUERY
```

```
100
```

```
1: in server:UPDATE 100 80
```

```
1 : commit:COMMIT
Going to commit
```

```
UPDATE
```

```
Sending:OK 80.00
```

```
1: in server:UPDATE 101 90
```

```
1 : commit:COMMIT
Going to commit
```

```
UPDATE
```

```
Sending:ERR Account 101 does not exist
```

```
smitha@shusmitha-VirtualBox: ~  
Write successful to server3:UPDATE 100 80  
Read successful from server1YES  
Read successfully from server2:  
YES  
REad successful from server3  
YES  
Write successful to server1  
COMMIT  
Write successful to server2  
COMMIT  
Write successful to server3  
COMMIT  
Read successful from server1 %s  
OK 80.00  
Read successful from server2 OK 80.00  
REad successful from server3  
OK 80.00  
Written final result to client  
Read successful from client:UPDATE 101 90  
Write successful to server1:UPDATE 101 90  
Write successful to server2  
UPDATE 101 90  
Write successful to server3:UPDATE 101 90  
Read successful from server1YES  
Read successfully from server2:  
YES  
REad successful from server3  
YES  
Write successful to server1  
COMMIT  
Write successful to server2  
COMMIT  
Write successful to server3  
COMMIT  
Read successful from server1 %s  
ERR Account 101 does not exist  
Read successful from server2 ERR Account 101 does not exist  
REad successful from server3  
ERR Account 101 does not exist  
Written final result to client
```

FRONTEND SERVER OUTPUT

```
Terminal File Edit View Search Terminal Help
shusmitha@shusmitha-VirtualBox:~$ ./clients localhost 8888
OK
Enter the command
Create 80
CREATE 80

Received : OK 100

QUERY 100
Enter the command
QUERY 100

Received : OK 80.00

UPDATE 100 80
Enter the command
UPDATE 100 80

Received : OK 80.00

UPDATE 101 90
Enter the command
UPDATE 101 90

Received : ERR Account 101 does not exist
```

OUTPUT OF CLIENT PROGRAM